


Logistic Regression

Logistic regression is a binary classification method based on logistic function.

Logistic function is a S-shape function whose output values are between 0 and 1 for any input values. It is defined as:

$$g(z) = \frac{1}{1 + e^{-z}}$$

A typical logistic function looks like this:



Logistic regression use a similar equation to predict an output value for any input value. A typical logistic regression equation is as follows:

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Here θ is a vector representing the weights or coefficient values related to x . Keep in mind that this equation is different from linear regression where a linear relationship exists in between the input values (x) and output values (y or $h_{\theta}(x)$), as in

$$h_{\theta}(x) = \theta^T x$$

The output value $h_{\theta}(x)$ denotes the probability of $y = 1$ on input x . That is,

$$h_{\theta}(x) = P(y = 1 | x, \theta)$$

This probability value is then transformed into a binary value (0 or 1) in order to make a probability prediction. Specifically, suppose we predict

$$y = 1 \text{ if } h_{\theta}(x) \geq 0.5$$

$$y = 0 \text{ if } h_{\theta}(x) < 0.5$$

That is equal to

$$y = 1 \text{ if } \theta^T x \geq 0$$

$$y = 0 \text{ if } \theta^T x < 0$$



Now it is clear that $\theta^T x = 0$ defines the **decision boundary** for the binary classification.

Similar to linear regression, we want to find the optimum θ that defines the decision boundary and mathematically we optimize the cost function defined as follows:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

Here the cost function for each (x, y) pair is defined in a way that when $y = 1$ and ideally $h_{\theta}(x) = 1$ or when $y = 0$ and ideally $h_{\theta}(x) = 0$, the cost function is minimum (0). Below figures shows how the cost function change with $h(x)$ for $y = 0$ and $y = 1$ respectively.




Similar to linear regression, the cost function $J(\theta)$ can be minimized using gradient descent. Below is the pseudo code:

repeat until converge:

$$\left\{ \begin{array}{l} \theta := \theta - \alpha / m * \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)} \end{array} \right\}$$

The figure below show how the cost function decrease with number of iterations for data in the Titanic Kaggle competition with slide modification (see Python code file).



References:

- 1, Machine Learning by Andrew Ng, from Coursera
- 2, Titanic Kaggle competition: <https://www.kaggle.com/c/titanic/data>