

Two Sigma Connect: Rental Listing Inquiries

<https://www.kaggle.com/c/two-sigma-connect-rental-listing-inquiries>

Cathy Qian, 2017

Outline

1. Problem description
2. Train Data analysis
3. Machine Learning Outline
4. Feature engineering
5. Machine learning Results Summary
6. Take-home message

1. Problem Description

- Predict the popularity of an apartment rental listing.
- Figure out key features responsible for the popularity of apartment rental listings.

training dataset:

9352 entries, 15 columns (with interest_level as target)

testing dataset:

74659 entries, 14 columns (without interest_level).

15 features

1. bathrooms: number of bathrooms, float
2. bedrooms: number of bedrooms, int
3. building_id: the id of the building, string
4. created: date and time when the post is created, string
5. description: description of the apartment, string
6. display_address: display address of the apartment in the posting, string
7. features: a list of features about this apartment, string
8. latitude: latitude of the apartment, float
9. listing_id: listing id of the apartment, int
10. longitude: longitude of the apartment, float
11. manager_id: id of the manager of the apartment, string
12. photos: a list of photo links. string
13. price: in USD, int
14. street_address: street address of the apartment, string
15. interest_level: this is the target variable. It has 3 categories: 'high', 'medium', 'low'

2. Train Data Analysis

Numerical features:

bathrooms: [0, 10], mean = 1.2

bedrooms:[0, 8.0], mean = 1.5

latitude:[0.000000, 44.883500]

longitude:[-118.271000, 0.000000]

listing_id:[6811957, 7753784], unique for each listing

price:[43, 4490000]

Texts:

created → extract day, hour, week of day

photos → extract number of photos

features → extract length of features or key words

description → extract length or key words

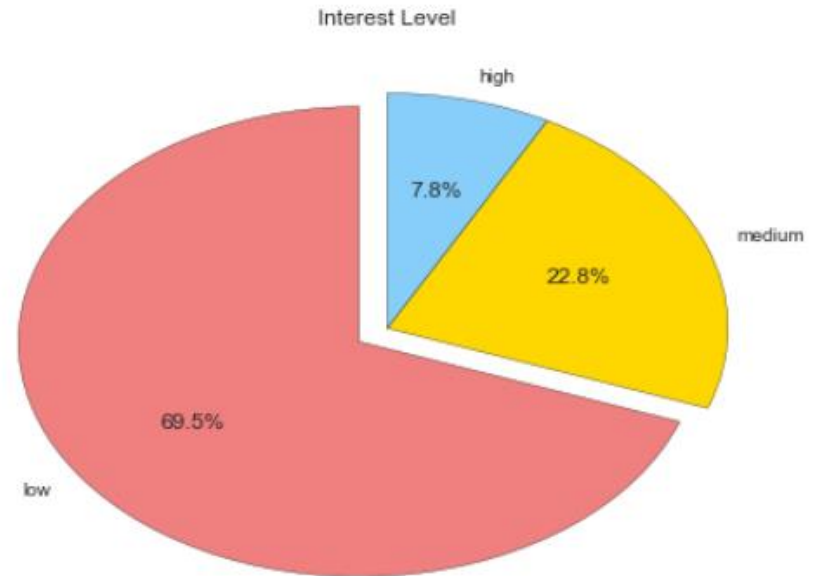
building_id: there're multiple listings with the same building_id

manager_id: there're multiple listings with the same manager_id

display_address: there're multiple listings with the same display_address

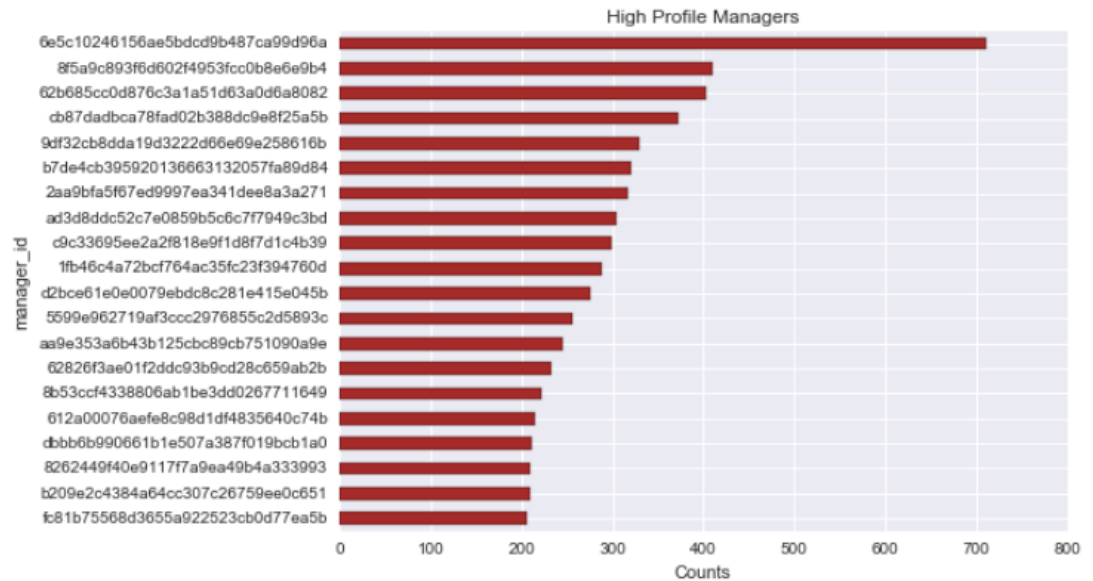
street_address: there're multiple listings with the same street_address

interest_level: distribution shown above

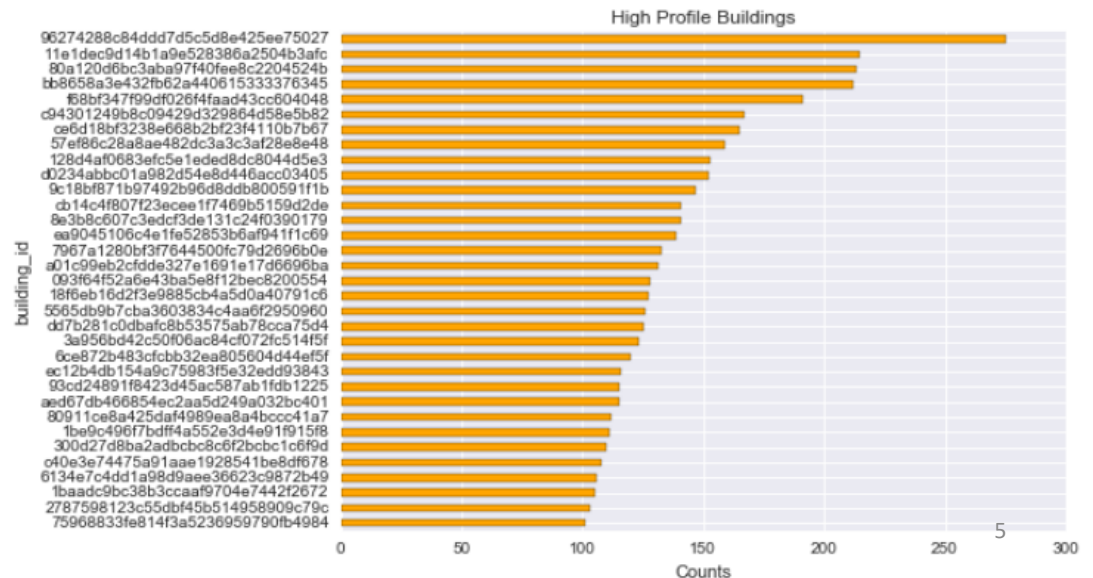


2. Train Data Analysis

manager_id with counts over 200



building_id with counts over 100

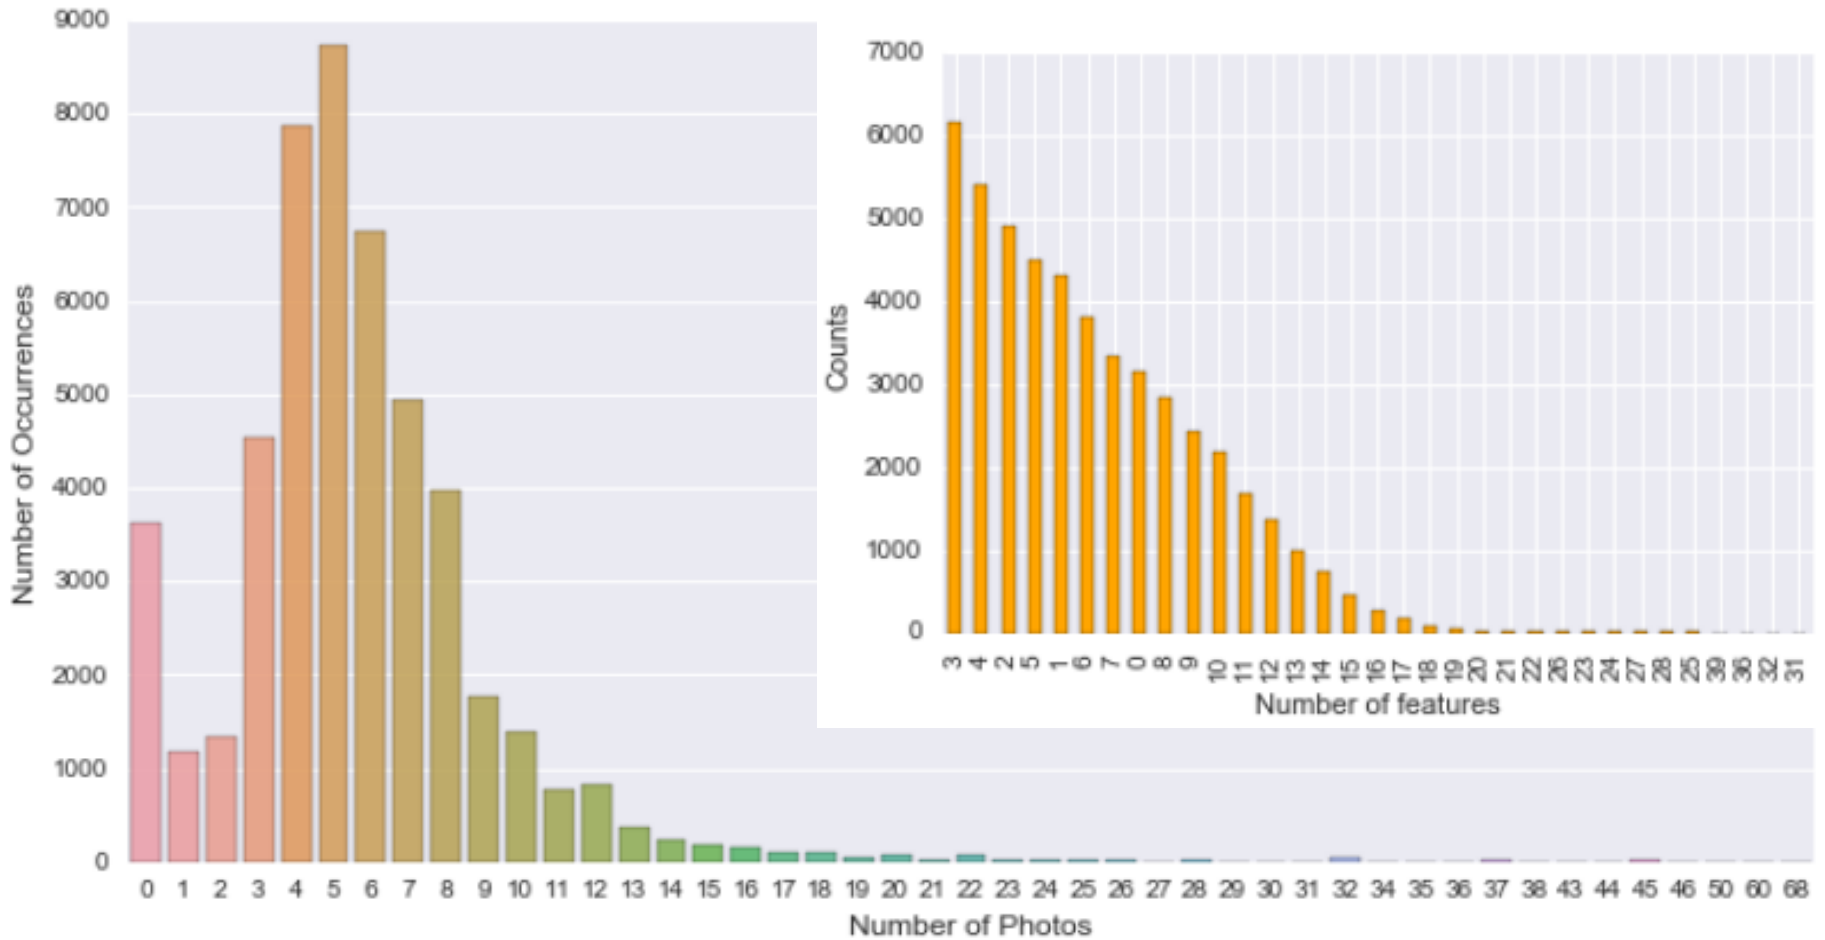


2. Train Data Analysis



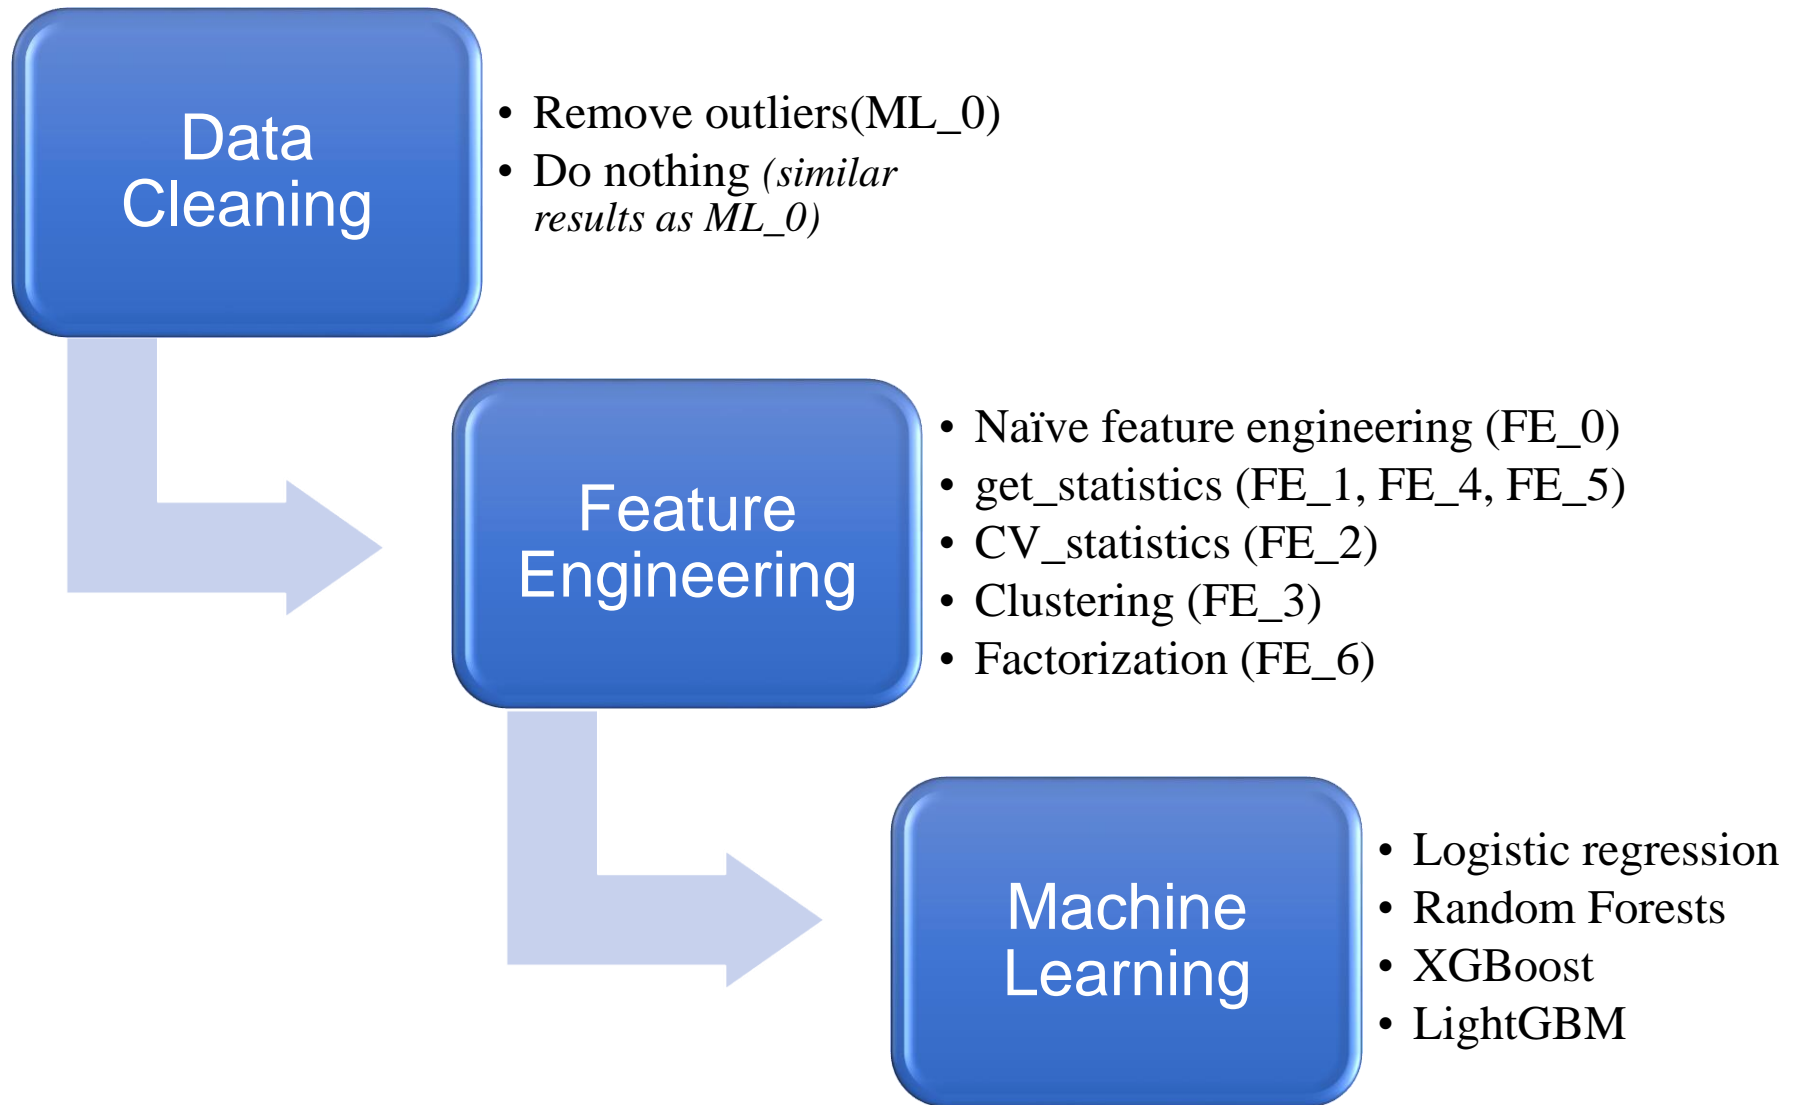
Only 99.5% data are shown.

2. Train Data Analysis



Skewed distribution

3. Machine Learning Outline



4. Feature Engineering

- Naïve feature engineering

- total number of rooms
- average price per room
- number of photos
- length of features
- number of words in description
- created day, month, hour

- get_statistics

group the dataframe by group column (manager_id, building_id), then calculate the count, mean, std, median, max, min of the target column (bathrooms, bedrooms, latitude, longitude, price etc) feature

- cv_statistics

calculate building_level = {manager_id: low_count, medium_count, high_count} Then update three new features: low_count%, medium_count% and high_count% for both train_df and test_df. If this manager_id only shows up in train_df but not test_df, nan is added.

encode categorical values into numerical values between 0 and n_classes - 1

- Clustering

categorize 'features' by the top ten features

separate Friday from the rest of the days and clustering the time of day into four categories

- Factorization

Encode input values as an enumerated type or categorical variable (i.e., 0, 1, 2,.....) and return the unique values

5. Machine learning Results Summary

Ref on logloss: <http://www.exegetic.biz/blog/2015/12/making-sense-logarithmic-loss/>

Log loss on X_validation/lb	Logistic Regression	Random Forests	XGBoost (without NAN)	XGBoost (withNAN)	LightGBM (withnan)	Notes
FE_0	0.70304 /0.72024	0.62085 /0.63383	0.631709 /0.62627	0.628339 /0.62627	0.6068221 /0.60777	Too simple FE
FE_1	0.667379 /1.00415	0.6153 /1.4447	0.572166 /1.15843	0.582721 /1.07354	0.550506 /1.16804	Overfitting?? Or data leaking
FE_2	0.459439 /test data contain NAN	0.4482 /test data contain NAN	0.408885 /0.87942	0.407328 /0.87324	0.40023587 /0.86359	Data leaking??
FE_3	0.795639/0.82395	0.8112 /0.82395	0.796590 /0.79797	0.796626 /0.796626	0.79183 /0.79195	Bad performance
FE_4	0.70040 /0.94479	0.61523 /1.35541	0.603335 /1.02034	0.600084 /0.94132	0.5619336 /1.17163	Overfitting?? Or data leaking
FE_5	N.A.	N.A.	N.A.	0.592260 /1.05636	0.534319 /1.27489	Overfitting?? Or data leaking
FE_6 (with listing_id)	0.646260 /0.64718	0.57552 /0.58216		0.550675 /0.55800	0.538097 /0.54073 0.53814(replace NAN with -1)/0.54114	Best result

6. Take-home message

Lessons learned from this competition:

- 1, Naïve Bayes and SVM performs really bad. Maybe because naïve assumption doesn't hold while SVM is good for “linear” separation which may not be the case here.
- 2, Merge test and train data before doing feature engineering may give better results than doing feature engineering solely on train data because this gives more data.
- 3, Tree based models like xgboost and lightGBM are not sensitive to features scales, so feature scaling is not needed.
- 4, LightGBM performs better and faster than xgboost in all investigated cases.
- 5, listing_id is important for getting small logloss value.
- 6, Among all tried ML algorithms, only XGBoost and lightGBM can handle NAN value.
- 7, More feature engineering, single algorithm training and ensembling/stacking are needed to decrease the logloss and increase the prediction accuracy given enough time. There is leaking in the dataset, which is subjected to further investigation.